

Implementation Note: Unrolling Parameters

Learning Algorithm

Have initial parameters $\theta^{(1)}, \theta^{(2)}, \theta^{(3)}$.

Unroll to get initial θ to pass to

`fminunc (@costFunction, initial θ , options)`

`function [jval, gradientVec] = costFunction(thetaVec)`

From `thetaVec`, get $\theta^{(1)}, \theta^{(2)}, \theta^{(3)}$.

Use forward prop/ back prop to compute $D^{(1)}, D^{(2)}, D^{(3)}$ and $J(\theta)$

Unroll $D^{(1)}, D^{(2)}, D^{(3)}$ to get `gradientVec`

Gradient Checking

- can approximate the derivative of cost f with:

$$\frac{\partial}{\partial \theta} J(\theta) \approx \frac{J(\theta - \epsilon) - J(\theta + \epsilon)}{2\epsilon}$$

Random Initialization

- 모든 θ 가중치를 0으로 초기화하는 것은 신경망에 부정확함
- Symmetry breaking

Putting It Together

Training a neural network

0. Pick a network architecture (뉴런 간 연결 패턴)
 - no. of input units, hidden layers, output units
 - Dimension of features $x^{(i)}$ (usually the more the better)
 - no. of classes
1. Randomly initialize weights θ
2. Implement forward propagation to get $h_{\theta}(x^{(i)})$ for any $x^{(i)}$
3. Implement code to compute cost function $J(\theta)$
4. Implement backprop to compute partial derivatives $\frac{\partial}{\partial \theta_{jk}} J(\theta)$
5. Use gradient checking
6. Use gradient descent or advanced optimization method with backpropagation to try to minimize $J(\theta)$ as a function of parameters θ
non-convex